

# Adaptive Particle Markov Chain Monte Carlo for Jump-Diffusion Models

## And a Shift to Differential Particle Filters

Michelle Ko

Supervised By  
Dr. Martin Lysy

University of Waterloo

May 29, 2023

# Table of Contents

Motivation and Mathematical Background

Particle Filtering

Particle MCMC

Differentiable Particle Filters

Discussion

References and Appendix

# Motivation

**Goal:** Recast asset price jump-diffusion as a state-space model to

- ▶ Recover the latent (unobserved) volatility, and
- ▶ Estimate model parameters,

based on observed asset price using a "particular" computational technique.

Example of a jump-diffusion model for asset price:

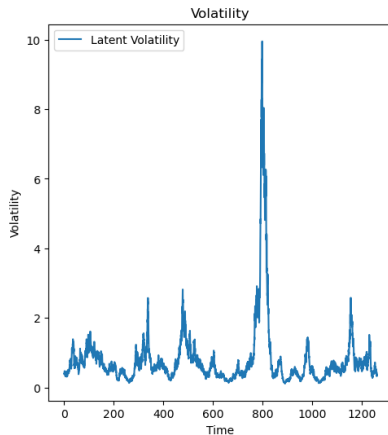
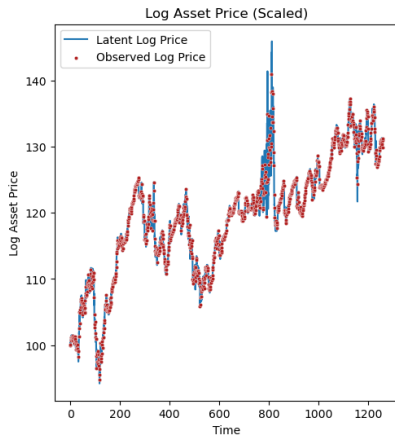
## Exponential Ornstein-Uhlenbeck (ExpOU)

Log asset price:  $dX_t = \alpha dt + \exp(Z_t)^{\frac{1}{2}} dW_t^x + V_t^x dN_t$

Log latent volatility:  $dZ_t = \kappa(\theta - Z_t)dt + \sigma dW_t^z + V_t^z dN_t$

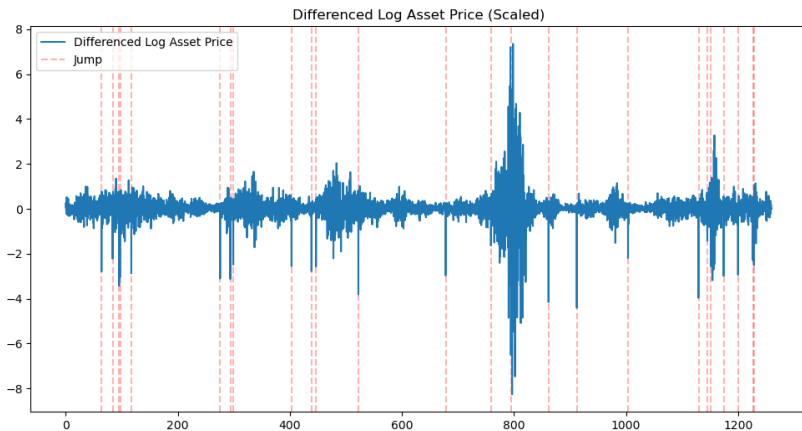
# Working Example

Consider a synthetic stock's daily price over 5 years that follows an ExpOU + Jump model.



## Working Example


The differenced log price gives a close approximation for daily return percentage.



# State-Space Representation

Volatility and log asset price,  $Z_t, X_t$ , are both continuous latent processes. Not a problem.

The observed log price,  $Y_t$ , is measured error-free  $\Rightarrow$  observation density is a Delta function. This becomes a problem later.



# Particle Filtering

Estimation method for the filtering distribution  $p(Z_{1:T} \mid X_{1:T}, \Theta)$   
[Gordon et al., 1993, Del Moral et al., 2001,  
Golightly and Wilkinson, 2008, Johannes et al., 2009].

1. **Propagate:** Sample from proposal  $Z_t^{(i)} \sim q(Z_t \mid \tilde{Z}_{t-1}, X_t, \Theta)$
2. **Re-weight:** Calculate incremental importance weights:

$$w_t^{(i)} = \frac{p_{\text{obs}}(X_t \mid Z_t^{(i)}, \Theta) p_{\text{trans}}(Z_t^{(i)} \mid \tilde{Z}_{t-1}^{(i)}, \Theta)}{q(Z_t^{(i)} \mid \tilde{Z}_{t-1}, X_t, \Theta)} \quad (1)$$

3. **Resample:** Resample particles with probability  $p_i \propto w_t^{(i)}$ :


$$\tilde{Z}_t \sim \text{Resampler}(Z_t, w_t^{(i)}) \quad (2)$$

**Importantly:** The marginal log-likelihood of  $\Theta$  given  $X_{1:T}$  "can" be unbiasedly estimated,

$$\hat{\ell} = \sum_{t=1}^T \log(\text{avg}_i(w_t^{(i)})) \quad (3)$$

# Particle Filtering Visualized


Prepare initial particles.






# Particle Filtering Visualized

Propagate particles with transition density.




# Particle Filtering Visualized

Re-weigh particles according to importance weights.




# Particle Filtering Visualized

Resample particles based on importance weights.



# Particle Filtering Visualized

Repeat!



# Transition Density: Discretizing Jump-Diffusions

Revisiting ExpOU with Jump...

- ▶  $m - 1$  inter-observations between each pair of observations
- ▶ Approximate the SDE using Euler-Maruyama scheme
- ▶ Use a Bernoulli distribution for jumps


Sampling from the transition density follows as [Golightly, 2009]:

1. Sample jump  $J \sim \text{Bernoulli}(\lambda \Delta t)$  and jump sizes  $V^x, V^z$
2. Sample  $(X_{i+1}, Z_{i+1})$  conditional on  $X_i, Z_i, J, V^x, V^z, \Theta$  using:

$$\pi(Z_{i+1} \mid Z_i, \Theta) = N(Z_i + \kappa(\theta - Z_i)\Delta t + V_{i+1}^z J_{i+1}, \sigma^2 \Delta t)$$

$$\pi(X_{i+1} \mid X_i, Z_i, \Theta) = N(X_i + \alpha \Delta t + V_{i+1}^x J_{i+1}, \exp(Z_i) \Delta t)$$

But what happens with no-error measurement?



# Proposal Density: Diffusion Bridge with Jumps

How to deal with this?

- Create a bridge conditioned on the next observation

$$Y_{t+1} = X_{t+1}$$

- This prevents all particle weights reducing to zero

The latent process in  $(t_j, t_M]$  is proposed [Golightly, 2009]: For  $i = j, \dots, M-1$ , first simulate  $M$  jumps and jump sizes. Then simulate  $Z_{i+1}^*$  recursively from its transition density. Then draw:

$$X_{i+1}^* \sim N\left(X_i^* + \frac{x_M - X_i^*}{M-i} + V_{i+1}^* J_{i+1}^* - \frac{\sum_{k=i+1}^M V_k^{x*} J_k^*}{M-i}, \frac{M-i-1}{M-i} Z_i^* \Delta t\right)$$

# Takeaway

We represent jump-diffusion model as a state-space model to:

- ▶ Use a particle filter.

We build a particle filter to:

- ▶ Recover the latent states given observation,
- ▶ Integrate over the latent states to obtain the marginal log-likelihood, which is almost always analytically intractable.

We construct a diffusion bridge to:

- ▶ Do the above even when observations are error-free.



# Particle MCMC

Particle Markov Chain Monte Carlo  $\Rightarrow$  When particle filtering is used within Markov chain Monte Carlo for drawing the posterior of the parameters [Andrieu et al., 2010]

- ▶ **Particle Marginal Metropolis-Hastings (PMMH):**

Metropolis-Hastings that uses the marginal likelihood estimated by particle filtering in the acceptance ratio

- ▶ **Particle Gibbs (PG):**

Alternate sampling between parameter and latent state using particle filters, i.e. draw parameter, draw a path based on parameter drawn, draw parameter based on path drawn...

Other options are Particle Metropolis-within-Gibbs, SMC<sup>2</sup>, etc.

# Adaptive PMCMC

PMCMC algorithms usually have tuning parameters:

- ▶ **PMMH:**

The step size  $\sigma_{rw}$  if the parameter proposal is a random walk

- ▶ **Particle Gibbs (PG):**

Dependent on how parameter conditioned on path is sampled

With methods proposed in [Roberts and Rosenthal, 2009],  
parameter tuning can be automatically done based on target  
acceptance rates, empirical covariance structure, etc.

# Preliminary Results

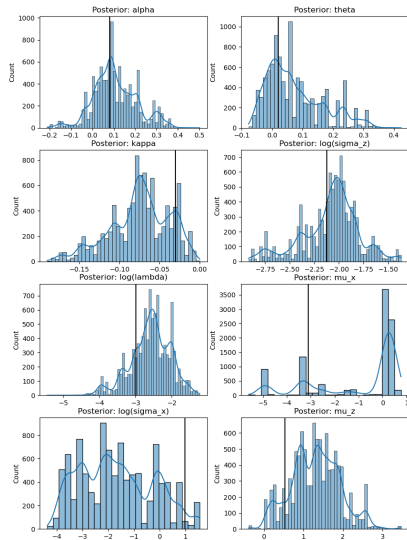
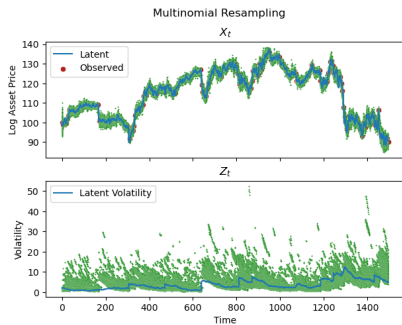
With a synthetic dataset generated by Heston + Jump model, Adaptive Particle Gibbs (APG) ran for around 9 minutes with:

- ▶ Number of observations: 300
- ▶ Resolution: 5
- ▶ Number of particles: 100
- ▶  $\Rightarrow$  150,000 operations per one marginal log-likelihood evaluation
- ▶ Number of MCMC iterations: 20,000

Obstacles:

- ▶ Extremely low acceptance rate on some parameters
- ▶ Computationally expensive to obtain plausible posterior

# Preliminary Results



# New Focus on Differentiable Particle Filters

## **Why do we want this?**

- ▶ If the PF-estimated marginal log-likelihood is differentiable, it turns into an optimization problem
- ▶ Gradient-based methods can be used, Autograd already well-implemented JAX, a high performance Python package [Bradbury et al., 2018]

## **How can we get this?**

1. Resampling method in the particle filter is differentiable
2. Random variables in the model have differentiable densities

# 1. Differentiable Resampler

Multinomial resampling is one of the traditional methods in particle filtering for the resampling step, i.e. choosing particle "ancestors":

- ▶  $Multinomial(\{w_t^{(i)}\})$
- ▶ Unbiased but NOT differentiable!


Consider instead a Gaussian approximation of the weighted particle distribution:

- ▶  $N(mean(\{Z_t^{(i)}\}, \{w_t^{(i)}\}), var(\{Z_t^{(i)}\}, \{w_t^{(i)}\}))$
- ▶ May be biased in some cases but differentiable!

# 1. Differentiable Resampler

When particles are resampled with Multinomial distribution:

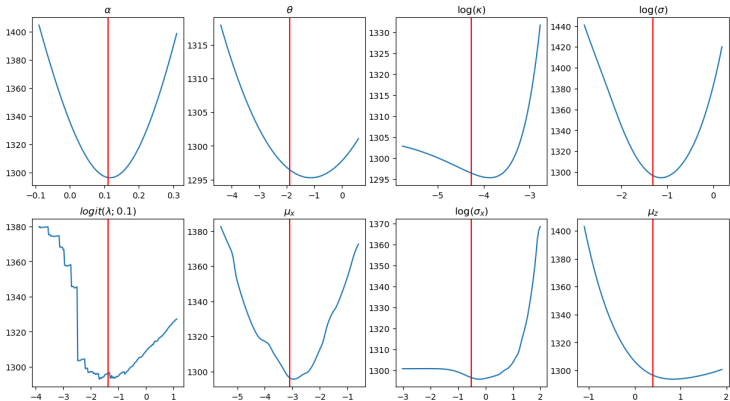
Marginal Negative Log Likelihood Projection: Multinomial Resampler



# 1. Differentiable Resampler

When particles are resampled with Gaussian approximation:

Marginal Negative Log Likelihood Projection: Gaussian Resampler







# 1. Differentiable Resampler

The two together:

Marginal Negative Log Likelihood Projection: Multinomial vs Gaussian Resampler



Why still jagged?



## 2. Differentiable Densities

Jump occurrence is a Bernoulli random variable:

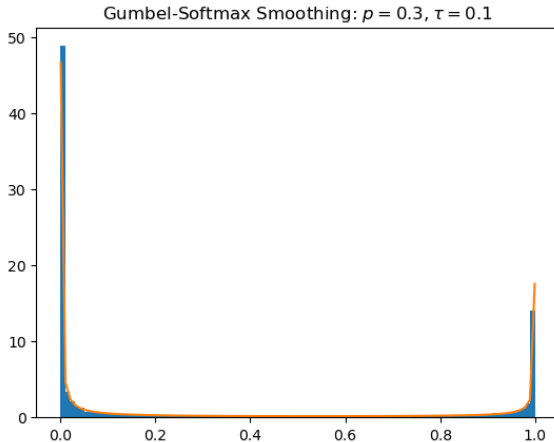
- ▶ Definitely NOT differentiable!
- ▶ Explains the jagged-ness of the marginal negative log-likelihood

We can employ a reparameterization trick:

- ▶ Gumbel-Softmax—comes with a tuning parameter  $\tau$  [Jang et al., 2017]
- ▶ As  $\tau \rightarrow 0$ , the distribution becomes Bernoulli
- ▶ Adds bias but differentiable!

## 2. Differentiable Densities

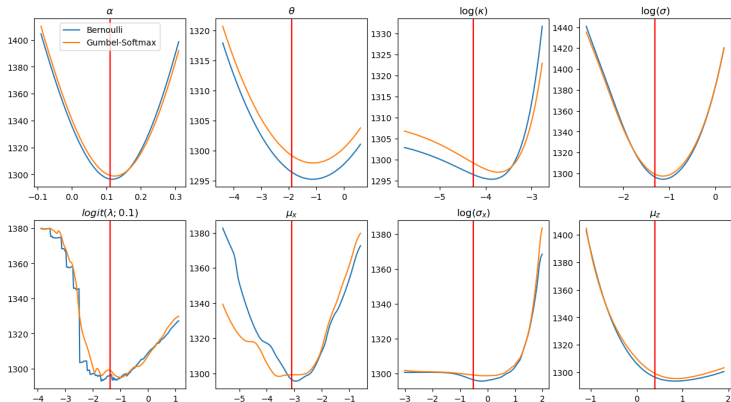
Distribution of  $(1 + \exp((L + \log \frac{1-p}{p})\tau^{-1}))^{-1}$ ,  $L \sim \text{logistic}(0, 1)$



## 2. Differentiable Densities



With  $\tau = 0.1$ , the minima is slightly shifted in some projection plots.

Marginal Negative Log Likelihood Projection: Bernoulli vs Gumbel-Softmax





## 2. Differentiable Densities

Increasing  $\tau$  means smoother marginal in  $\lambda$ .



## 2. Differentiable Densities

Bias is prominent in  $\alpha$  as  $\tau$  increases.



## Results

With a slight modification of  $\lambda$ , Gradient Descent ran for around 10 minutes with:

- ▶ Number of observations:  $252 \times 3 = 756$
- ▶ Resolution: 5
- ▶ Number of particles: 200
- ▶  $\Rightarrow$  756,000 operations for one marginal log-likelihood evaluation

Parameter	$\alpha$	$\theta$	$\kappa$	$\sigma$
True	0.11	-1.9	0.014	0.27
Estimated	0.18	-2.75	0.017	0.22

Parameter	$\lambda$	$\mu_x$	$\sigma_x$	$\mu_z$
True	0.050	-3.1	0.60	0.64
Estimated	0.063	-3.75	1.43	0.59



# Filtered Latent States

With true parameters:



# Filtered Latent States

With parameter estimates obtained from Reparameterized model:



# Discussion

Cases that did not work well:

- ▶ Small  $\lambda$ : Jumps are too rare
- ▶ Small  $\tau$ : Not enough smoothing
- ▶ Small  $\mu_x, \mu_z$ : Jump sizes are negligent
- ▶ Other factors: Wrong model specification, "bad seed", etc.

Areas of further investigation:

- ▶ Tradeoff between  $\tau$  and bias: Bias correction?
- ▶ Distribution of  $\hat{\Theta}$ : Run optimizer with different seeds?
- ▶ Real life data: Daily SP 500 data?
- ▶ Correlated processes: Between  $X_t, Z_t$  or between jump sizes?
- ▶ Portfolio: Multiple assets that follow jump diffusion?

# Acknowledgement

The PFJAX team:

- ▶ Martin Lysy
- ▶ Jonathan Ramkissoo
- ▶ Pranav Subramani
- ▶ Mohan Wu
- ▶ Kanika Chopra

# References I



Andrieu, C., Doucet, A., and Holenstein, R. (2010).

Particle markov chain monte carlo methods.

*Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 72(3):269–342.



Bradbury, J., Frostig, R., Hawkins, P., Johnson, M., Leary, C., Maclaurin, D., Necula, G., Paszke, A., VanderPlas, J., Wanderman-Milne, S., and Zhang, Q. (2018).

Jax: composable transformations of Python+NumPy programs.



Del Moral, P., Jacod, J., and Protter, P. (2001).

The monte-carlo method for filtering with discrete-time observations.

*Probability Theory and Related Fields*, 120:346–368.

# References II



Golightly, A. (2009).

Bayesian filtering for jump-diffusions with application to stochastic volatility.

*Journal of Computational and Graphical Statistics*,  
18:384–400.



Golightly, A. and Wilkinson, D. (2008).

Bayesian inference for nonlinear multivariate diffusion models observed with error.

*Computational Statistics and Data Analysis*, 52:1674–1693.



Gordon, N., Salmond, D., and Smith, A. (1993).

Novel approach to nonlinear/non-gaussian bayesian state estimation.

*IEE Proceedings-F*, 140:107–113.



Jang, E., Gu, S., and Poole, B. (2017).

Categorical reparameterization with gumbel-softmax.

## References III



Johannes, M., Polson, N., and Stroud, J. (2009).

Optimal filtering of jump diffusions: Extracting latent states from asset prices.

*Review of Financial Studies*, 22:2759–2799.



Roberts, G. and Rosenthal, J. (2009).


Examples of adaptive mcmc.

*Journal of Computational and Graphical Statistics*, 18:349–367.

# Background

State-space model is specified by:

- ▶ Latent state with transition density:  $p_{\text{trans}}(Z_t \mid Z_{t-1}, \Theta)$
- ▶ Observation density:  $p_{\text{obs}}(X_t \mid Z_t, \Theta)$



We are interested in the estimation of the model parameter  $\Theta$  and latent state  $Z_{1:T}$ .